

'Always and Everywhere': Compute is the New Inflation

Yahali Pessach

Over the past five years, inflation was the foremost macroeconomic constraint facing the United States. But over the next five years, inflation will play a diminishing role. As artificial intelligence expands productive capacity, the US will enter an era of abundance in goods and services. Access to compute will become the new scarcity shaping economic outcomes. Capital flows into technology and strategic industries will outpace realized productivity gains, placing pressure on the infrastructure supporting this expansion. The disruption will accelerate labor displacement and encourage modern mercantilism, reinforcing state intervention through 2031. This analysis focuses on the United States, where the compute build-out and mercantilist shift are moving fastest, but the dynamics will spread elsewhere.

Forecasts

A continued rise in capital expenditure on AI impacts businesses and consumers alike.

- F1 A 67% chance the four largest US tech firms' capex hits \$800bn in 2027.
- F2 A 66% chance North American data-center vacancy stays below 2% to 2031.

In turn, a rise in costs of production...

- F3 A 52% chance US commercial power in Q4 2028 is 12%+ above Q4 2024.
- F4 A 52% chance US retail power in Q4 2030 is 23%+ above Q4 2024.

A rise in American mercantilism yields a rise in protectionism and state involvement.

- F5 A 67% chance the average US tariff stays over 6% of imports to 2031.
- F6 A 70% chance the US government holds stakes in 58+ firms by 2028.

In turn, rivals escalate...

- F7 A 39% chance China's US Treasury holdings fall below \$500bn by 2029.
- F8 A 53% chance China's magnet exports to the US drop 40% vs 2024 by 2028.

A convergence of AI integration and rising mercantilism leads to a structurally higher price of money.

- F9 A 64% chance US core CPI inflation averages 2.7%+ every year to 2031.
- F10 An 84% chance the fed funds rate never falls below 2.5% before 2032.
- F11 A 70% chance the 10-year yield averages 4.5%+ over 2027 to 2031.
- F12 A 57% chance central banks buy 800+ tons of gold in 2027 and 2028.

Capital and work reorganize faster than productivity materializes, so protection, once in place, endures.

- F13 A 47% chance an S&P 500 firm takes a \$5bn+ AI write-down by 2031.
- F14 An 80% chance US office-and-admin jobs in May 2031 are 1.5%+ below 2025.
- F15 A 68% chance US productivity growth averages under 1.75% a year to 2031.
- F16 An 80% chance the US tariff falls no more than a point below its 2027 level until 2031.

Framework and synthesis


Forecasts are generated from interconnected networks of causes and consequences in a model for an open US economy. In short, as AI expands capacity in ordinary goods, the constraint on the American economy moves from labor to compute. In turn, the inputs to a compute build-out grow scarce, with power the first among them, and that scarcity draws in the state. Once the state steps in, the cost of its response in restricting free trade leads to a structurally higher price of money. This doesn't necessarily affect prosperity, but it does prompt a reinforcing cycle of intervention and investment over coming years.

Forces Driving the New Economic Order

Bets on compute mount in the price of power.

Microsoft, Alphabet, Amazon and Meta spent \$413bn in 2025, with another \$720bn expected in 2026. This comes as investment in computing across the wider economy is accelerating with them. Because the magnificent seven treat this as a race they cannot afford to lose, their spending has become highly price inelastic. Since 2018 American electricity prices have risen steadily as natural gas, which sets most of the wholesale price, has fallen. This is one reflection of prices being primarily demand-driven, even as supply bottlenecks come and go.

US power prices and natural gas, index (Jan 2018 = 100)




US Construction and Data-Center Power Are Rising Together




Scarcity leads to intervention, then scarcity again.

Faced with shortages of strategic inputs, the American bureaucracy has not allowed AI-adjacent markets to clear. Instead, it has intensified tariffs and export controls. The response has taken two forms, deeper state involvement in strategic industries and a more restrictive trade regime. Federal investment stakes in strategic firms have risen. Meanwhile, duties alone now exceed a tenth of goods import values, the highest share in more than sixty years. As protection intended to sustain the compute build-out raises costs, those costs become embedded in a higher price level. Because the build-out resembles a strategic game of chicken, it cannot easily be paused. Demand remains elevated while supply stays constrained, turning compute into a persistent source of inflation rather than a passing shock.

US tariff rate and duty revenue




Tariffs Have Been Layered On




Prices get muddy as trade is redrawn.

A more closed America has pushed global trade to reroute. Imports from China have almost halved since 2022, while imports from Vietnam and Mexico have surged. Much of this reflects Chinese production redirected through third countries, creating a divergence in inflation dynamics across the globe. As America absorbs higher costs from protectionism, China faces deflationary pressure from excess industrial capacity redirected away from the US market.

US goods imports by source, index (2017 = 100)




Consumer prices, % change on a year earlier



Dollar debased?


Weakening confidence in the dollar as a neutral reserve asset has increasingly pushed central banks to gold. This shift reflects a broader erosion of trust in the existing monetary order, contributing to stickier inflation expectations, while a rising term premium has kept long-term yields elevated despite Fed easing.

Gold Is Climbing Even as Real Yields Climb, a Break From History



This suggests that monetary policy is becoming less powerful. As structural and geopolitical risks account for a larger share of long-term yields, central banks have less influence over financial conditions, making inflation harder to contain without imposing greater costs on growth.

Change in the US Treasury yield curve since January 2026



Deriving Assumptions on the Future of Work and Trade

Focusing on the United States, I treat Chinese state actions as exogenous shocks to the American economy. I do this for two reasons. First, the compute build-out associated with AI development has been concentrated overwhelmingly in the United States and China, with foreign actors years away from comparable development. Second, the shift towards modern mercantilism in the Western sphere has largely been driven by the United States, whereas in the Eastern sphere it was already more established.

To model this trajectory, I begin with six premises on the evolution of state intervention and technological progress. I assign each premise a value drawn from a normal distribution centered on the estimates below, with feasible ranges defining 90% intervals. I then run a Monte Carlo simulation year by year from 2027 to 2031 within a stochastic decision tree. Across 200,000 simulations, each forecast probability represents the share of paths in which that outcome occurs. I chose this approach specifically to capture feedback across economic and political dynamics that emerge endogenously within the model. Take:

M1. For every 10% by which demand for a strategic input exceeds available supply, the state adds roughly 1.1 percentage points of tariff within a year.

M2. For every two points of new American tariff, China's answer costs the build-out about one percentage point of capacity growth the following year, aimed where its advantage is greatest.

M3. Of the protection in place at the start of a year, around 95% remains by the end of it, assuming no measures are repealed.

C1. Spending on compute grows at around 30% a year, half the pace of the past two years, but still enough to double expenditure by 2029.

C2. Deliverable capacity grows by around 25% a year, with the compute build-out limited by its slowest constraint.

C3. The compute build-out adds roughly 0.8 percentage points to productivity growth, but only after adoption catches up with investment.


Initial Outputs

This yields preliminary estimates on a variety of metrics on the economy:

Median simulated world	2027	2028	2029	2030	2031
tariff collected, % of goods imports (F5, F16)	6.62	6.70	6.70	6.70	6.70
federal equity stakes, number held (F6)	44	59	74	88	103
data-center vacancy, % (F2)	1.28	1.27	1.27	1.27	1.27
big-four capital spending, \$bn (F1)	904	1,091	1,282	1,482	1,712
business electricity, % above 2024 (F3)	8.78	12.27	15.89	19.81	24.02
household electricity, % above 2024 (F4)	12.47	16.07	19.81	23.87	28.22
core inflation, % (F9)	2.93	2.84	2.92	2.94	2.94
federal funds rate, % (F10)	3.82	3.55	3.79	3.83	3.84
ten-year Treasury yield, % (F11)	4.66	4.39	4.63	4.67	4.68
central-bank gold, tons (F12)	823	834	834	834	833
largest single-firm write-down, \$bn (F13)	0.00	0.00	0.69	0.72	0.89
clerical employment, % below 2025 (F14)	0.60	1.31	2.10	2.95	3.84
productivity growth, % a year (F15)	1.44	1.51	1.57	1.61	1.64

A Bird's Eye View of the Model

The United States open-economy model is formulated as a nonlinear recursive system driven by six stochastic inputs, which represent the model's primary sources of uncertainty. These inputs propagate through a deterministic computational graph, with each intermediate variable generated by fixed functional relationships calibrated from historical data. Because the system is recursive, outcomes become state variables, allowing shocks to accumulate over time through endogenous feedback. The resulting forecasts are jointly determined, with their covariance structure emerging from stochastic drivers and recursive evolution of the system.




forecast, and the quantity it asks about	bar	10th	50th	90th	chance	the 200,000 worlds, and the bar on them
F1 capex in 2027, \$bn	≥ 800	637	904	1,192	67%	
F2 vacancy, highest year, %	< 2	0.94	1.28	7.14	66%	
F3 commercial power 2028, % vs 2024	≥ 12	11.72	12.27	20.35	52%	
F4 retail power 2030, % vs 2024	≥ 23	22.10	23.87	49.21	52%	
F5 tariff, lowest year, %	≥ 6	5.11	6.39	6.79	67%	
F6 state stakes in 2028	≥ 58	56.64	58.98	65.80	70%	
F7 China's US Treasuries, \$bn	< 500	441	504	504	39%	
F8 magnets to US, % below 2024	≥ 40	30.94	40.54	49.46	53%	
F9 core, lowest year, %	≥ 2.7	2.46	2.76	2.90	64%	
F10 funds, lowest year, %	≥ 2.5	2.11	3.35	3.72	84%	
F11 ten-year, average, %	≥ 4.5	4.33	4.62	6.12	70%	
F12 gold, lower year, tonnes	≥ 800	739	808	849	57%	
F13 largest write-down, \$bn	≥ 5	0.00	2.89	72.27	47%	
F14 clerical fall by 2031, %	≥ 1.5	0.48	3.84	7.48	80%	
F15 productivity, average, %	< 1.75	1.11	1.56	2.05	68%	
F16 tariff, worst fall from 2027, pts	≥ -1	-1.35	-0.07	0.00	80%	

Premises on Technological Progress

C1. Spending on compute grows at around 30% a year, half the pace of the past two years, but still enough to double expenditure by 2029 (F1, F13).


Here, I assume that compute investment remains elevated, but that the recent pace of expansion moderates. Microsoft, Alphabet, Amazon and Meta spent \$250bn on capital investment in 2024 and \$413bn in 2025. They expect to spend around \$720bn in 2026, implying growth of roughly 67% a year for two years running.⁷ I do not expect this pace to continue indefinitely, however. Sustaining it to 2029 would push spending by the four firms alone above \$3.2 trillion a year, close to a tenth of American output, without a clear revenue base to support it. A 30% annual growth rate instead gets them to around \$1.5 trillion, still remarkable but easier to justify. I do not expect a sharper slowdown because none of the four can easily stop investing while the others continue, and because the spending is largely being funded by profits and strong balance sheets rather than the debt that carried the telecoms boom.⁸

The main uncertainty is whether expectations adjust faster than the technology generates returns. If AI revenues disappoint, investment could slow sharply. However, I think this risk is overstated for the AI capital expenditure build-out. The largest firms can fund it from existing cash flows, and much of the infrastructure behind it will remain valuable beyond any single AI product cycle. The chips themselves are more vulnerable because they depreciate within a few years, creating the possibility of a large write-down even if the broader AI build-out continues (F13). I place the range for compute investment growth between -20% and +80% relative to the baseline assumption, with the lower bound reflecting the possibility of a correction like the fall in telecoms investment after 2000.




C2. Deliverable capacity grows by around 25% a year, with the compute build-out limited by its slowest constraint (F2, F3, F4).

Here, I postulate that deliverable compute capacity grows rapidly but remains constrained by the slowest-moving input. North American data-center capacity was a third higher in early 2026 than it had been just a year earlier.⁹ However, the ceiling on delivery is set by the scarcest input. At present, that constraint is power, which is becoming harder to secure before new supply arrives: gas turbines are sold out into 2031, while large transformers take around four years from order to delivery.¹⁰ This is Liebig's law of the minimum in practice. Capital can accelerate investment, but it cannot remove physical lead times. One qualification offered by Nathan Rosenberg is that bottlenecks often redirect innovation towards the constraint itself, because that is where returns are highest.¹¹ I therefore assume capacity growth of around 25% a year, with a feasible range of 5% to 45%. The upper end depends on whether today's power constraints can be resolved by the end of the decade.




Source: LBNL "Queued Up" 2024 (queue, end-2023); EIA (installed capacity), by resource type




Source: CBRE, PJM, Wood Mackenzie, manufacturer disclosures (years from order to delivery)

C3. The compute build-out adds around 0.8 percentage points to productivity growth, but only after adoption catches up with investment (F14, F15).

Here, I expect the productivity impact of AI to arrive with a lag, as firms take time to reorganize around the technology. Early investment therefore does not imply an immediate rise in measured output per hour. Paul David's history of the dynamo shows why the payoff from a new technology can arrive late: electricity only transformed American factories after firms took time to reorganize production around it.¹² Brynjolfsson, Rock and Syverson formalize this idea through their J-curve, where investment precedes productivity because complementary changes take time.¹³ I therefore estimate that AI eventually adds around 0.8 percentage points to annual productivity growth, though only part of that effect is likely to appear before the end of the decade. I place the feasible range at roughly 0 to 1.7 percentage points, with the lower bound reflecting delayed adoption and the upside tempered by Robert Gordon's argument that recent technologies may prove less transformative than those of the past.¹⁴



Source: BLS output per hour, nonfarm business, via FRED; 2027-31 is the model's median forecast



Source: BEA real fixed investment, BLS output per hour, via FRED


Premises on Modern Mercantilism

M1. For every 10% by which demand for a strategic input exceeds available supply, the state adds roughly 1.1 percentage points of tariff within a year (F5, F16).

Here, I estimate the relationship between supply shortages and tariff escalation. In 2017, America collected 1.63% in duties on imported goods. By early 2026, that figure had risen to 10.15%, with almost 90% of the increase occurring in 2025 and 2026 alone.¹ Tariffs have tended to accumulate over time, as new measures have been layered onto existing schedules rather than replacing earlier ones. Averaged across the eight years since the first Section 301 actions in 2018, this represents just over 1 percentage point of additional tariff per year. This period coincided with persistent shortages in strategic inputs, as demand exceeded available supply by roughly 10%. Transformer and turbine backlogs lengthened, while rare-earth magnet constraints intensified. Rather than extrapolate the exceptional pace of 2025–26 indefinitely, I smooth the recent surge across the full period. Relative to the estimated supply shortfall, this implies roughly 11 percentage points of additional tariff for a full unit of excess demand, which is the parameter used in the model. The range, from 3 percentage points if the recent acceleration proves temporary to 19 percentage points if it marks a new era of mercantilism, captures both scenarios.


I expect this relationship to persist because the forces behind it are unlikely to reverse quickly. Critical technologies and resources are increasingly viewed not only as economic inputs but as strategic assets, making governments less willing to tolerate dependence on foreign suppliers. This shift has support

across both parties and extends beyond any single administration. It also reflects a recurring historical pattern: Alexander Hamilton and Friedrich List argued that strategic industries may justify higher costs when foreign dependence creates vulnerability.² Karl Polanyi described the same political reflex from another angle: when markets disrupt societies faster than they can adjust, pressure builds for state protection.³ The era of economic liberalization from the 1990s to the 2010s may therefore prove to be an exception rather than the long-run equilibrium. The main uncertainty is whether shortages produce tariffs or alternative interventions such as subsidies and exemptions, as occurred during the chip shortage. I assume tariffs remain the fastest and most visible trade response, but where policymakers choose other tools, M1 will cause the model to overstate subsequent reciprocal measures.



M2. For every two points of new American tariff, China's answer costs the build-out about one percentage point of capacity growth the following year, aimed where its advantage is greatest (F7, F8).


Here, I capture the idea that trade dependence creates leverage beyond economic exchange. This follows Albert Hirschman's argument that the country better able to withstand disruption holds influence over the one that cannot.⁴ China's position in refined rare earths and magnets illustrates this mechanism. Rather than matching tariffs directly, it has targeted areas where its supply-chain advantage is greatest, including gallium and germanium controls in July 2023 and magnet licensing with extraterritorial reach in April 2025.⁵ Each followed an American measure, suggesting that strategic dependencies will be used as geopolitical leverage in future. China's ability to retaliate remains asymmetric, though. Restricting exports risks damaging its own manufacturing base and export markets, while targeted controls on critical inputs can impose disproportionate costs abroad. M2 therefore assumes that Chinese retaliation is roughly half the size of the initial American tariff shock, concentrated where its supply advantage is greatest. This logic also has limits, so I place the feasible range for Chinese retaliation relative to the initial American action at 0.15 to 0.85.



M3. Of the protection in place at the start of a year, around 95% remains by the end of it, assuming no measures are repealed (F9 to F12).

Here, I incorporate that tariffs have proved sticky. Once introduced, they acquire a domestic constituency and carry no expiry date. Between 2019 and 2024, when almost no major new tariffs were introduced, the rate above the pre-2018 base fell only from 1.47 percentage points to 0.93, an annual decline of 8.7%.¹ The erosion came through exclusions and changing trade patterns rather than repeal, so the Section 301 duties largely survived a change of administration. Even the Supreme Court's ruling against


the IEEPA tariffs in Learning Resources v. Trump removed less protection than it seemed to, because much of the schedule stands on other authorities.⁶ I set persistence at 95% rather than the measured 91.3%, reflecting that past decay came through exclusions that have since lapsed, while future rollbacks are more likely to target broad consumer-facing tariffs than measures on semiconductors and other critical inputs. I estimated a feasible range of 0.90 to 1.00, reflecting that tariffs are largely controlled by the executive in the US and that reciprocal tariffs often retain public support elsewhere.



Methodology

Define the six assumptions. The model begins with six assumptions, each expressed as a central estimate and a plausible range. The range reflects uncertainty and is interpreted as containing 90% of possible outcomes. Values are also bounded where necessary, to eliminate any impossible outcomes.

Simulate outcomes. The model draws values for the six assumptions and runs the economy forward from 2027 to 2031. It repeats this process 200,000 times to generate a distribution of outcomes. Two of the underlying uncertainties are discrete rather than continuous, so they are modeled through four weighted cases: (1) whether investment delivers productivity gains within five years and (2) whether existing protection remains in place. Each simulation first selects one of four cases: a 25% chance that capacity expands and protection increases, a 20% chance that capacity expands and protection fades, a 45% chance that expansion is delayed while protection increases, and a 10% chance that both the compute build-out and the shift toward modern mercantilism weaken. I set ‘delay’ as likelier than ‘delivery’ (55% against 45%) because the J-curve behind C3 suggests adoption will lag investment. I also set protection to persist in 70% of worlds because that is what the stickiness argued in M1 and M3 implies. Finally, escalation is likelier when the build-out fades (82% against 56%), because a disappointing boom leaves the state facing more dislocation.



Run the model forward. Once the six assumptions are set, the model applies the same equations year by year from 2027 to 2031 using observed data and derived quantities.

Calculate probabilities. Each forecast is defined by a threshold which I picked, with the probability as the share of simulations in which that threshold is reached.


Test sensitivity. I rerun the model with each assumption moved across its full range to assess dependence on individual inputs. A few modelling choices require clarification. First, the model does not include a separate recession mechanism because downturns are difficult to parameterize ex ante and are instead captured through the lower bounds of C1 and C3. Second, tariffs are allowed to affect inflation beyond a one-off price-level adjustment because recent experience suggests pass-through can be staggered through exclusions and supply-chain adjustments. Third, the ten-year yield is linked to the funds-rate path using the current curve slope, leaving term-premium variation outside the model. Finally, non-tariff core inflation is anchored at 2% to separate persistent inflation from the tariff and power-price channels being modelled. Each assumption is then tested through the counterfactual scenarios.

Counterfactuals

Each counterfactual examines what would happen if one of the six initial conditions is wrong. To do this, I move one of the six conditions from the bottom to the top of its plausible range, while leaving the other five unchanged. I rerun the same 200,000 simulations and then recalculate all sixteen forecast probabilities.

How far each forecast moves when one assumption is set at the high bound of its range as opposed to the lower bound

	M1	M2	M3	C1	C2	C3
F1	-	-	-	+85	-	-
F2	-	-	-	+84	-37	-
F3	-	-	+1	+87	-40	-
F4	-	-	+2	+88	-40	-
F5	+22	-	+51	+54	-25	+5
F6	+20	-	+62	+40	-21	+1
F7	+28	-	+14	+82	-40	+1
F8	-	+90	-	-	-	-
F9	+32	-	+55	-32	+20	+2
F10	+19	-	+6	-57	+31	-
F11	+26	-	+55	+40	-17	+3
F12	+10	-	+85	+32	-16	+2
F13	-3	-1	-1	-54	+32	-
F14	-	-	-	-	+5	+75
F15	-	-	-	-	-17	-85
F16	+21	-	+47	+35	-16	+5



The left-hand panel varies each assumption in turn and shows the percentage-point change in each forecast. Capital expenditure, data-center construction and electricity prices are driven mainly by compute demand and capacity growth. Employment and productivity depend on the timing of productivity gains. Inflation, rates, yields and central-bank gold purchases depend on the persistence of protection. China's role is narrow, affecting only rare-earth export controls due to limited data and a broader modelling limitation. The model captures China raising build-out costs but not the deflationary effect of redirected overcapacity through third countries. If anything, this omission biases inflation forecasts upward.

The right-hand panel highlights four asymmetries. If protection unwinds, the probability of sustained central-bank gold buying falls from 57% to 9%, as reserve demand for gold follows geopolitical fragmentation. If the investment boom fades, inflation risk rises because tariffs remain embedded while the loss of capacity growth removes a source of future disinflation. If China escalates, rare-earth export controls become far likelier, yet the broader outlook barely moves because the model is driven mainly by domestic investment. If productivity arrives sooner, the probability of sub-1.75% productivity growth falls sharply, but protection persists because it reflects a political response rather than an economic constraint.

References

1. Bureau of Economic Analysis, customs duties and imports of goods, via FRED (B235RC1Q027SBEA, A255RC1Q027SBEA).
2. Alexander Hamilton, Report on the Subject of Manufactures (1791); Friedrich List, The National System of Political Economy (1841).
3. Karl Polanyi, The Great Transformation (1944).
4. Albert O. Hirschman, National Power and the Structure of Foreign Trade (1945).
5. Chinese Ministry of Commerce announcements, July 2023, December 2024, April 2025; suspension November 2025.
6. Supreme Court of the United States, Learning Resources v. Trump, February 2026.
7. Company filings and earnings calls: Microsoft, Alphabet, Amazon, Meta, 2024 to 2026.
8. Carlota Perez, Technological Revolutions and Financial Capital (2002).
9. CBRE, Global Data Center Trends 2026, first-quarter data.
10. Turbine and transformer lead times as reported by manufacturers; PJM interconnection queue disclosures.
11. Nathan Rosenberg, The Direction of Technological Change: Inducement Mechanisms and Focusing Devices (1969).
12. Bureau of Labor Statistics, nonfarm business output per hour, via FRED (OPHNFB).
13. Paul A. David, The Dynamo and the Computer (American Economic Review, 1990).
14. Erik Brynjolfsson, Daniel Rock and Chad Syverson, The Productivity J-Curve (American Economic Journal: Macroeconomics, 2021).
15. Robert J. Gordon, The Rise and Fall of American Growth (2016).
16. Ray Dalio, Principles for Navigating Big Debt Crises (2018); How Countries Go Broke (2025).

Data Sources

Customs duties	fred.stlouisfed.org/series/B235RC1Q027SBEA
Goods imports	fred.stlouisfed.org/series/A255RC1Q027SBEA
Core CPI	fred.stlouisfed.org/series/CPILFESL
CPI electricity	fred.stlouisfed.org/series/CUSR0000SEHF01
PPI commercial power	fred.stlouisfed.org/series/WPU0542
Output per hour	fred.stlouisfed.org/series/OPHNFB
Federal funds rate	fred.stlouisfed.org/series/FEDFUNDS
10-year Treasury	fred.stlouisfed.org/series/GS10
GDP	fred.stlouisfed.org/series/GDP
Consumption	fred.stlouisfed.org/series/PCE
Data centers	cbre.com/insights/books/global-data-center-trends-2026
Capex, Microsoft	microsoft.com/en-us/investor
Capex, Alphabet	abc.xyz/investor
Capex, Amazon	ir.aboutamazon.com
Capex, Meta	investor.atmeta.com
Clerical employment	bls.gov/oes (OEWS, occupation 43-0000)
Central-bank gold	gold.org/goldhub/data
Tariff cash and accrual	fiscal.treasury.gov/reports-statements/mts
Power in business costs	bea.gov/industry/input-output-accounts-data
Data-center load	eia.gov/electricity

Code Base

Available at github.com/yahalipessach/Bridgewater-Forecasting-the-Future.